

Contents

Report: adrian-research/minerva-rl-benchmark-6	1
Summary	1
Ant [opt=adam, adv=gae]	1
Best-Run GIFs	2
HalfCheetah [opt=adam, adv=gae]	2
Best-Run GIFs	2
HalfCheetah [opt=adam, adv=returns]	2
Best-Run GIFs	2
Humanoid [opt=adam, adv=gae]	2
Best-Run GIFs	2
Humanoid [opt=adam, adv=returns]	2
Best-Run GIFs	3
Walker2d [opt=adam, adv=gae]	3
Best-Run GIFs	3
Walker2d [opt=adam, adv=returns]	3
Best-Run GIFs	3
cheetah/run [opt=adam, adv=none]	3
Best-Run GIFs	3
humanoid/walk [opt=adam, adv=none]	3
Best-Run GIFs	4
pendulum [opt=unknown, adv=unknown]	4
Best-Run GIFs	4

Report: adrian-research/minerva-rl-benchmark-6

- Generated: 2026-02-15 13:05:15
- Included runs: 12 (`_step > 1`)
- Algorithm key source: run config `command`
- Environment key source: run config `env` + `optimizer_type` + `advantage_estimator`
- Metric: `eval/return_mean`

No plottable metric history was found for the selected runs.

Summary

Environment / Optimizer / Adv Type	Algorithms	Runs
Ant [opt=adam, adv=gae]	ppo	1
HalfCheetah [opt=adam, adv=gae]	ppo	1
HalfCheetah [opt=adam, adv=returns]	vmpo	1
Humanoid [opt=adam, adv=gae]	ppo	1
Humanoid [opt=adam, adv=returns]	vmpo	1
Walker2d [opt=adam, adv=gae]	ppo	1
Walker2d [opt=adam, adv=returns]	vmpo	1
cheetah/run [opt=adam, adv=none]	mpo	1
humanoid/walk [opt=adam, adv=none]	mpo	1
pendulum [opt=unknown, adv=unknown]	pendulum-vmpo-gtrxl-seed0, r2d2-gtrxl, vmpo-gtrxl	3

Ant [opt=adam, adv=gae]

Run	Algorithm	_step	eval/return_mean
ppo_Ant-v5-adam-gae_20260215-122440	ppo	499712	287.203

Best-Run GIFs

Best run per algorithm by `eval/return_max` (fallback `eval/return_mean`).

Algorithm	Run	Best Metric	Preview	Status
ppo	ppo_Ant-v5-adam-gae_20260215-122440	964.610	-	skipped: Checkpoint file not found.

HalfCheetah [opt=adam, adv=gae]

Run	Algorithm	_step	eval/return_mean
ppo_HalfCheetah-v5-adam-gae_20260215-101843	ppo	499712	1735.166

Best-Run GIFs

Best run per algorithm by `eval/return_max` (fallback `eval/return_mean`).

Algorithm	Run	Best Metric	Preview	Status
ppo	ppo_HalfCheetah-v5-adam-gae_20260215-101843	1816.913	-	skipped: Checkpoint file not found.

HalfCheetah [opt=adam, adv=returns]

Run	Algorithm	_step	eval/return_mean
vmpo_HalfCheetah-v5-adam>Returns_20260215-101844	vmpo	3000000	840.225

Best-Run GIFs

Best run per algorithm by `eval/return_max` (fallback `eval/return_mean`).

Algorithm	Run	Best Metric	Preview	Status
vmpo	vmpo_HalfCheetah-v5-adam-returns_20260215-101844	1746.357	-	skipped: Checkpoint file not found.

Humanoid [opt=adam, adv=gae]

Run	Algorithm	_step	eval/return_mean
ppo_Humanoid-v5-adam-gae_20260215-115235	ppo	499712	564.861

Best-Run GIFs

Best run per algorithm by `eval/return_max` (fallback `eval/return_mean`).

Algorithm	Run	Best Metric	Preview	Status
ppo	ppo_Humanoid-v5-adam-gae_20260215-115235	905.211	-	skipped: Checkpoint file not found.

Humanoid [opt=adam, adv=returns]

Run	Algorithm	_step	eval/return_mean
vmpo_Humanoid-v5-adam>Returns_20260215-120644	vmpo	2392171	569.381

Best-Run GIFs

Best run per algorithm by `eval/return_max` (fallback `eval/return_mean`).

Algorithm	Run	Best Metric	Preview	Status
vmpo	vmpo_Humanoid-v5-adam>Returns_20260215-120644	989.586	-	skipped: Checkpoint file not found.

Walker2d [opt=adam, adv=gae]

Run	Algorithm	_step	eval/return_mean
ppo_Walker2d-v5-adam-gae_20260215-112104	ppo	499712	626.997

Best-Run GIFs

Best run per algorithm by `eval/return_max` (fallback `eval/return_mean`).

Algorithm	Run	Best Metric	Preview	Status
ppo	ppo_Walker2d-v5-adam-gae_20260215-112104	734.302	-	skipped: Checkpoint file not found.

Walker2d [opt=adam, adv=returns]

Run	Algorithm	_step	eval/return_mean
vmpo_Walker2d-v5-adam>Returns_20260215-111111	vmpo	3000000	631.832

Best-Run GIFs

Best run per algorithm by `eval/return_max` (fallback `eval/return_mean`).

Algorithm	Run	Best Metric	Preview	Status
vmpo	vmpo_Walker2d-v5-adam>Returns_20260215-111111	1165.683	-	skipped: Checkpoint file not found.

cheetah/run [opt=adam, adv=none]

Run	Algorithm	_step	eval/return_mean
mpo_dm_control-cheetah-run-adam-none_20260215-101849	mpo	400000	525.603

Best-Run GIFs

Best run per algorithm by `eval/return_max` (fallback `eval/return_mean`).

Algorithm	Run	Best Metric	Preview	Status
mpo	mpo_dm_control-cheetah-run-adam-none_20260215-101849	530.872	-	skipped: Checkpoint file not

humanoid/walk [opt=adam, adv=none]

Run	Algorithm	_step	eval/return_mean
mpo_dm_control-humanoid-walk-adam-none_20260215-112009	mpo	140242	1.312

Best-Run GIFs

Best run per algorithm by `eval/return_max` (fallback `eval/return_mean`).

Algorithm	Run	Best Metric	Preview	Status
mpo	mpo_dm_control-humanoid-walk-adam-none_20260215-112009	11.224	-	skipped: Checkpoint file r

pendulum [opt=unknown, adv=unknown]

Run	Algorithm	_step	eval/return_mean
pendulum-vmpo-gtrxl-seed0_pendulum_20260215-130137	pendulum-vmpo-gtrxl-seed0	2009	-1390.144
r2d2-gtrxl_pendulum_20260215-130409	r2d2-gtrxl	21	-1713.304
vmpo-gtrxl_pendulum_20260215-130240	vmpo-gtrxl	331	-1307.331

Best-Run GIFs

Best run per algorithm by `eval/return_max` (fallback `eval/return_mean`).

Algorithm	Run	Best Metric	Preview	Status
pendulum-vmpo-gtrxl-seed0	pendulum-vmpo-gtrxl-seed0_pendulum_20260215-130137	-1029.172	-	skipped: Ch
r2d2-gtrxl	r2d2-gtrxl_pendulum_20260215-130409	-1575.469	-	skipped: Ch
vmpo-gtrxl	vmpo-gtrxl_pendulum_20260215-130240	-1014.560	-	skipped: Ch